



Artificial Intelligence
Prof. Björn Ommer
HCI & IWR

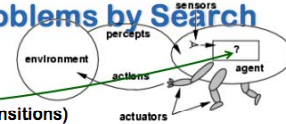





Nov 07, 2017 - Informed Search

Review: Solving Problems by Search

- Given:**
 - State space
 - possible actions (= state transitions)
- Goal:**
 - reach a specific state (or set of states w/ certain characteristics)
- Solution:**
 - sequence of actions (leading from initial state to goal)
- Find solution by searching for optimal seq. of actions**



State

5	4	
6	1	8
7	3	2

Node

parent, action

depth = 6

g = 6

state

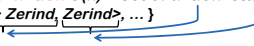
B. Ommer | ommer@uni-heidelberg.de

Single-state problem formulation

A **problem** is defined by four items:

- initial state** e.g., "at Arad"
- actions** or **successor function** $S(x)$ = set of action-state pairs
e.g., $S(\text{Arad}) = \{ \langle \text{Arad} \rightarrow \text{Zerind}, \text{Zerind} \rangle, \dots \}$
- goal test**, can be
 - explicit**, e.g., $x = \text{"at Bucharest"}$
 - implicit**, e.g., $\text{NoDirt}(x)$
- path cost** (additive)
 - e.g., sum of distances, number of actions executed, etc.
 - $c(x, a, y)$ is the **step cost**, assumed to be ≥ 0

A **solution** is a sequence of actions leading from the initial state to a goal state



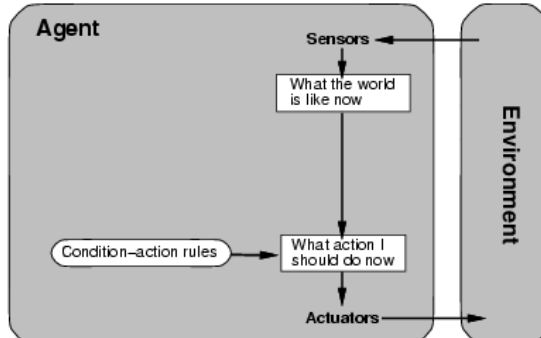
B. Ommer | ommer@uni-heidelberg.de

Review: Solving Problems by Search

- Uninformed Search:**
 - only local info about present state
 - e.g., "in goal state?", "how much has agent earned up till now?"
- Informed Search:**
 - Estimate of distance to goal from present state
- Desirability of**
 - present state as is vs. based on potential successors of present state

B. Ommer | ommer@uni-heidelberg.de

Search-based Agent



B. Ommer | ommer@uni-heidelberg.de

Outline – Informed search algorithms

- Best-first search
- Greedy best-first search
- A* search
- Heuristics
- Local search algorithms
- Hill-climbing search
- Simulated annealing search
- Local beam search
- Genetic algorithms

B. Ommer | ommer@uni-heidelberg.de

Review: Tree search

```

function TREE-SEARCH(problem, fringe) returns a solution, or failure
  fringe ← INSERT(MAKE-NODE(INITIAL-STATE(problem)), fringe)
  loop do
    if fringe is empty then return failure
    node ← REMOVE-FRONT(fringe)
    if GOAL-TEST[problem] applied to STATE(node) succeeds return node
    fringe ← INSERTALL(EXPAND(node, problem), fringe)

```

- A search strategy is defined by picking the **order of node expansion**


B. Ommer | ommer@uni-heidelberg.de

Best-first search

- Idea: use an **evaluation function** $f(n)$ for each node
 - estimate of "desirability"
 - Expand most desirable unexpanded node
- Implementation:
 - Order the nodes in fringe in decreasing order of desirability
- Special cases:
 - greedy best-first search
 - A* search

B. Ommer | ommer@uni-heidelberg.de

Romania with step costs in km



B. Ommer | ommer@uni-heidelberg.de

Greedy best-first search

- Evaluation function $h(n)$ (heuristic)
 - = estimate of cost from n to goal
- e.g., $h_{SLD}(n)$ = straight-line distance from n to Bucharest
- Greedy best-first search expands the node that **appears** to be closest to goal

B. Ommer | ommer@uni-heidelberg.de

Greedy best-first search example




B. Ommer | ommer@uni-heidelberg.de

Greedy best-first search example



B. Ommer | ommer@uni-heidelberg.de

Greedy best-first search example



B. Ommer | ommer@uni-heidelberg.de

Greedy best-first search example



B. Ommer | ommer@uni-heidelberg.de

Properties of greedy best-first search

- **Complete?** No – can get stuck in loops, e.g., $A \rightarrow B \rightarrow A \rightarrow B \rightarrow \dots$
- **Time?** $O(b^m)$, but a good heuristic can give dramatic improvement
- **Space?** $O(b^m)$ -- keeps all nodes in memory
- **Optimal?** No



B. Ommer | ommer@uni-heidelberg.de

A* search

- Idea: avoid expanding paths that are already expensive
- Evaluation function $f(n) = g(n) + h(n)$
 - $g(n)$ = cost so far to reach n
 - $h(n)$ = estimated cost from n to goal
 - $f(n)$ = estimated total cost of path through n to goal


B. Ommer | ommer@uni-heidelberg.de

A* search example




B. Ommer | ommer@uni-heidelberg.de

A* search example




B. Ommer | ommer@uni-heidelberg.de

A* search example



B. Ommer | ommer@uni-heidelberg.de

A* search example




B. Ommer | ommer@uni-heidelberg.de

A* search example



B. Ommer | ommer@uni-heidelberg.de

A* search example



B. Ommer | ommer@uni-heidelberg.de

Admissible heuristics


- A heuristic $h(n)$ is **admissible** if for every node n , $h(n) \leq h^*(n)$, where $h^*(n)$ is the **true** cost to reach the goal state from n .
- An admissible heuristic **never overestimates** the cost to reach the goal, i.e., it is **optimistic**.
- Example: $h_{SLD}(n)$ (never overestimates the actual road distance)
- Theorem:** If $h(n)$ is admissible, A* using TREE-SEARCH is optimal

B. Ommer | ommer@uni-heidelberg.de

Consistent heuristics

- A heuristic is **consistent** if for every node n , every successor n' of n generated by any action a , $h(n) \leq c(n, a, n') + h(n')$.
- If h is consistent, we have


$$\begin{aligned} f(n') &= g(n') + h(n') \\ &= g(n) + c(n, a, n') + h(n') \\ &\geq g(n) + h(n) \\ &= f(n) \end{aligned}$$
- i.e., $f(n)$ is non-decreasing along any p
- Theorem:** If $h(n)$ is consistent, A* using GRAPH-SEARCH is optimal



B. Ommer | ommer@uni-heidelberg.de

Optimality of A* (proof)

- Suppose some suboptimal goal G_2 has been generated and is in the queue. Let n be an unexpanded node in the queue such that n is on a shortest path to an optimal goal G .




- $f(G_2) = g(G_2)$ since $h(G_2) = 0$
- $g(G_2) > g(G)$ since G_2 is suboptimal
- $f(G) = g(G)$ since $h(G) = 0$
- $f(G_2) > f(G)$ from above

B. Ommer | ommer@uni-heidelberg.de

Optimality of A* (proof)

- Suppose some suboptimal goal G_2 has been generated and is in the queue. Let n be an unexpanded node in the queue such that n is on a shortest path to an optimal goal G .




- $f(G_2) > f(G)$ from above
 - $h(n) \leq h^*(n)$ since h is admissible
 - $g(n) + h(n) \leq g(n) + h^*(n)$
 - $f(n) \leq f(G)$
- Hence $f(G_2) > f(n)$, and A* will never select G_2 for expansion

B. Ommer | ommer@uni-heidelberg.de

Optimality of A*

- A* expands nodes in order of increasing f value
- Gradually adds "f-contours" of nodes
- Contour i has all nodes with $f=f_i$, where $f_i < f_{i+1}$



B. Ommer | ommer@uni-heidelberg.de

Properties of A*

- Complete?** Yes (unless there are infinitely many nodes with $f \leq f(G)$)
- Time?** Exponential [relative error in h * length of soln.]
- Space?** Keeps all nodes in memory
- Optimal?** Yes: cannot expand f_{i+1} until f_i is finished

- A* expands all nodes with $f(n) < C^*$
- A* expands some nodes with $f(n) = C^*$
- A* expands no nodes with $f(n) > C^*$


...Now: improve speed by better heuristics

B. Ommer | ommer@uni-heidelberg.de

Admissible heuristics

E.g., for the 8-puzzle:

- $h_1(n)$ = number of misplaced tiles
- $h_2(n)$ = total Manhattan distance (i.e., no. of squares from desired location of each tile)




- $h_1(S) = ?$
- $h_2(S) = ?$

B. Ommer | ommer@uni-heidelberg.de

Admissible heuristics

E.g., for the 8-puzzle:

- $h_1(n)$ = number of misplaced tiles
- $h_2(n)$ = total Manhattan distance (i.e., no. of squares from desired location of each tile)



- $h_1(S) = ?$ 8
- $h_2(S) = ?$ $3+1+2+2+2+3+3+2 = 18$

B. Ommer | ommer@uni-heidelberg.de

Dominance

- If $h_2(n) \geq h_1(n)$ for all n (both admissible) then h_2 **dominates** $h_1 \Rightarrow h_2$ is better for search
- Typical search costs (average number of nodes expanded):
 - $d=12$ IDS = 3,644,035 nodes
A*(h_1) = 227 nodes
A*(h_2) = 73 nodes
 - $d=24$ IDS = too many nodes
A*(h_1) = 39,135 nodes
A*(h_2) = 1,641 nodes
- Given any admissible heuristics h_a, h_b ,
 $h(n) = \max(h_a(n); h_b(n))$
is also admissible and dominates h_a, h_b

B. Ommer | ommer@uni-heidelberg.de

Relaxed Problems

- A problem with fewer restrictions on the actions is called a **relaxed problem**
- Removing constraints (relaxation) creates extra edges to search graph
 \Rightarrow state space graph of relaxed problem is **supergraph** of original one
- Extra edges: optimal solution in original problem is also solution of relaxed problem
- But relaxed problem may have better solutions, since extra edges can provide shortcuts

B. Ommer | ommer@uni-heidelberg.de


Relaxed Problems

- The cost of an optimal solution to a relaxed problem is an **admissible heuristic** for the original problem
 - If the rules of the 8-puzzle are relaxed so that a tile can move **anywhere**, then $h_1(n)$ gives the shortest solution
 - If the rules are relaxed so that a tile can move to **any adjacent square**, then $h_2(n)$ gives the shortest solution
- Key point: the optimal solution cost of a relaxed problem is not greater than the optimal solution cost of the real problem

B. Ommer | ommer@uni-heidelberg.de

Relaxed Problems contd.

- Well-known example: travelling salesperson problem (TSP)
- Find the shortest tour visiting all cities exactly once




- Minimum spanning tree can be computed in $O(n^2)$ and is a lower bound on the shortest (open = exactly one visit) tour

B. Ommer | ommer@uni-heidelberg.de

Example: TSP

- Start with any complete tour, perform pairwise exchanges



- Variants of this approach get within 1% of optimal very quickly with thousands of cities

B. Ommer | ommer@uni-heidelberg.de

Summary

- Heuristic functions estimate costs of shortest paths
- Good heuristics can dramatically reduce search cost
- Greedy best-first search expands lowest h
 - incomplete and not always optimal
- A* search expands lowest $g + h$
 - complete and optimal
 - also optimally efficient (up to tie-breaks, for forward search)
- Admissible heuristics can be derived from exact solution of relaxed problems

B. Ommer | ommer@uni-heidelberg.de

Outline – Informed search algorithms

- Best-first search
- Greedy best-first search
- A* search
- Heuristics
- Local search algorithms
- Hill-climbing search
- Simulated annealing search
- Local beam search
- Genetic algorithms

B. Ommer | ommer@uni-heidelberg.de

Local search algorithms


- In many optimization problems, the **path** to the goal is irrelevant; the goal state itself is the solution
- State space = set of "complete" configurations
- Find configuration satisfying constraints, e.g., n-queens
- In such cases, we can use **local search algorithms**
- keep a single "current" state, try to improve it

B. Ommer | ommer@uni-heidelberg.de

Example: n-queens

- Put n queens on an $n \times n$ board with no two queens on the same row, column, or diagonal

- Move a queen to reduce number of conflicts




- Almost always solves n-queens problems almost instantaneously for very large n , e.g., $n=1$ million

B. Ommer | ommer@uni-heidelberg.de

Hill-climbing search (gradient ascend)

- "Like climbing Everest in thick fog with amnesia"




```

function HILL-CLIMBING(problem) returns a state that is a local maximum
inputs: problem, a problem
local variables: current, a node
                 neighbor, a node
current ← MAKE-NODE(INITIAL-STATE[problem])
loop do
  neighbor ← a highest-valued successor of current
  if VALUE[neighbor] ≤ VALUE[current] then return STATE[current]
  current ← neighbor
  
```

B. Ommer | ommer@uni-heidelberg.de

Hill-climbing search


- Problem: depending on initial state, can get stuck in local maxima



- Random-restart hill climbing overcomes local maxima; trivially complete
- Random sideways moves escape from shoulders loop on at maxima

B. Ommer | ommer@uni-heidelberg.de

Hill-climbing search: 8-queens problem



- h = number of pairs of queens that are attacking each other, either directly or indirectly
- A local minimum with $h = 1$

B. Ommer | ommer@uni-heidelberg.de


Simulated annealing search

- Most minimization strategies find the *nearest local minimum*
- Standard strategy:
 - Generate trial point based on current estimates
 - Evaluate function at proposed location
 - Accept new value if it improves solution
- Idea: escape local maxima by allowing some "bad" moves but **gradually decrease** their frequency

B. Oettermann | oettermann@uni-heidelberg.de

Solution

- We need a strategy to find other minima
- This means, we must sometimes select new points that do not improve solution



- How?

B. Oettermann | oettermann@uni-heidelberg.de

Annealing

- One manner in which crystals are formed
- Gradual cooling of liquid ...
 - At high temperatures, molecules move freely
 - At low temperatures, molecules are "stuck"
- If cooling is slow
 - Low energy, organized crystal lattice formed



B. Oettermann | oettermann@uni-heidelberg.de

Simulated Annealing

- Analogy with thermodynamics
- Incorporate a temperature parameter into the minimization procedure
- At high temperatures, explore parameter space
- At lower temperatures, restrict exploration
- Consider decreasing series of temperatures
- For each temperature, iterate these steps:
 - Propose an update and evaluate function
 - Accept updates that improve solution
 - Accept some updates that don't improve solution
 - Acceptance probability depends on "temperature" parameter

B. Oettermann | oettermann@uni-heidelberg.de

Example Application - TSP

- The traveling salesman problem
 - Salesman must visit every city in a set
 - Given distances between pairs of cities
 - Find the shortest route through the set
- No practical deterministic algorithms for finding optimal solution are known...
 - ... simulated annealing and other stochastic methods can do quite well

B. Oettermann | oettermann@uni-heidelberg.de

Update Scheme

- A good scheme should be able to:
 - Connect any two possible paths
 - Propose improvements to good solutions
- Some possible update schemes:
 - Swap a pair of cities in current path
 - Invert a segment in current path

B. Oettermann | oettermann@uni-heidelberg.de

Example Trajectories in State Space



B. Ommer | ommer@uni-heidelberg.de

Updating


- **Metropolis (1953), Hastings (1970)**
 - Define a set of conditions that, if met, ensure the random walk will sample from probability distribution at equilibrium
 - In theory
- Recommendations apply to how changes are proposed and accepted

B. Ommer | ommer@uni-heidelberg.de

Accepting an Update

- The **Metropolis criterion**: Change from E_0 to E with probability

$$\min\left(1, \exp\left\{-\frac{(E - E_0)}{T}\right\}\right)$$




- Given sufficient time, leads to equilibrium state

B. Ommer | ommer@uni-heidelberg.de

Simulated Annealing and Gradient Descent

- Perform gradient descent on original cost fct + random noise (noise energy prop. to temp T)



B. Ommer | ommer@uni-heidelberg.de

Key Requirement: Irreducibility


- **All states must communicate**
 - We can get everywhere from any arbitrary starting point
 - Starting point should not affect results
- If Q is matrix of proposal probabilities
 - Either $Q_{ij} > 0$ for all possible states i and j
 - Some integer P exists where $(Q^P)_{ij} > 0$ for all i, j
 - After P steps we can reach any j from any i

B. Ommer | ommer@uni-heidelberg.de

- Probability of state with energy k is

$$P(E = k) \propto \exp\left(-\frac{k}{T}\right)$$

- At low T , probability is concentrated in low energy states



Simulated Annealing Recipe

1. Select starting temperature and initial parameter values
2. Randomly select a new point in the neighborhood of the original
3. Compare the two points using the *Metropolis criterion*
4. Repeat steps 2 and 3 until system reaches equilibrium state...
 - In practice, repeat the process N times for large N
5. Decrease temperature and repeat the above steps, stop when system reaches frozen state

B. Ommer | ommer@uni-heidelberg.de

Simulated annealing search

- Idea: escape local maxima by allowing some "bad" moves but **gradually decrease** their frequency

```

function SIMULATED-ANNEALING( problem, schedule) returns a solution state
inputs:  problem, a problem
        schedule, a mapping from time to "temperature"
local variables: current, a node
               next, a node
               T, a "temperature" controlling prob. of downward steps
current ← MAKE-NODE(INITIAL-STATE[problem])
for t ← 1 to ∞ do
  T ← schedule[t]
  if T = 0 then return current
  next ← a randomly selected successor of current
  ΔE ← VALUE[next] - VALUE[current]
  if ΔE > 0 then current ← next
  else current ← next only with probability e-ΔE/T

```

B. Ommer | ommer@uni-heidelberg.de

Practical Issues

- The maximum temperature
- Scheme for decreasing temperature
 - Cooling schedules
- Strategy for proposing updates
 - E.g.: Select a component to update
 - & Sample from within plausible range

B. Ommer | ommer@uni-heidelberg.de

Properties of simulated annealing search

- At fixed "temperature" T, state occupation probability reaches Boltzmann distribution

$$p(x) = \frac{e^{-\frac{E(x)}{T}}}{\sum e^{-\frac{E(x)}{T}}}$$
- T decreased slowly enough \Rightarrow always reach best state x^* because

$$\frac{e^{-\frac{E(x^*)}{T}}}{e^{-\frac{E(x)}{T}}} = e^{\frac{E(x^*) - E(x)}{T}} \gg 1 \text{ for small } T$$
- However, annealing can be exponentially slow
- Devised by Metropolis et al., 1953, for physical process modelling
- Widely used in VLSI layout, airline scheduling, etc

B. Ommer | ommer@uni-heidelberg.de

Local beam search

- Keep track of k states rather than just 1; choose top k of all their successors
 - Start with k randomly generated states
 - At each iteration, all the successors of all k states are generated
 - Not the same as k searches run in parallel! Searches that find good states recruit other searches to join them
 - If any one is a goal state, stop; else select the k best successors from the complete list and repeat.
- Observe the close analogy to natural selection!

B. Ommer | ommer@uni-heidelberg.de


Genetic algorithms

- A successor state is generated by combining two parent states
- Start with k randomly generated states (**population**)
- A state is represented as a string over a finite alphabet (often a string of 0s and 1s)
- Evaluation function (**fitness function**). Higher values for better states.
- Produce the next generation of states by selection, crossover, and mutation
- If there is more info on cost fct, more effective optimization strategies are typically preferred

B. Ommer | ommer@uni-heidelberg.de

Genetic algorithms

= stochastic local beam search + generate successors from pairs of states



Fitness Selection Pairs Cross-Over Mutation


- Fitness function: number of non-attacking pairs of queens (min = 0, max = $8 \times 7/2 = 28$)
- $24/(24+23+20+11) = 31\%$
- $23/(24+23+20+11) = 29\%$ etc

B. Ommer | ommer@uni-heidelberg.de

Genetic algorithms

GAs require states encoded as strings (GPs use programs)

Crossover helps iff substrings are meaningful components



GAs \neq evolution: e.g., real genes encode replication machinery!

B. Ommer | ommer@uni-heidelberg.de