




Artificial Intelligence
Prof. Björn Ommer
HCI & IWR




RUPRECHT-KARLS-
UNIVERSITÄT
HEIDELBERG




Nov 14, 2017 - CSP



Outline – Constraint Satisfaction Problems

- Constraint Satisfaction Problems (CSP)
- Backtracking search for CSPs
- Local search for CSPs


B. Ommer | ommer@uni-heidelberg.de




Constraint satisfaction problems (CSPs)

- **Standard search problem:**
 - **state** is a "black box" – any data structure that supports successor function, heuristic function, and goal test
- **CSP:**
 - **state** is defined by **variables** X_i with **values** from **domain** D_i
 - **goal test** is a set of **constraints** specifying allowable combinations of values for subsets of variables
- Simple example of a **formal representation language**
- Allows useful **general-purpose** algorithms with more power than standard search algorithms

B. Ommer | ommer@uni-heidelberg.de




Example: Map-Coloring




- **Variables** WA, NT, Q, NSW, V, SA, T
- **Domains** $D_i = \{red, green, blue\}$
- **Constraints:** adjacent regions must have different colors
 - e.g., $WA \neq NT$, or $(WA, NT) \in \{(red, green), (red, blue), (green, red), (green, blue), (blue, red), (blue, green)\}$

B. Ommer | ommer@uni-heidelberg.de




Example: Map-Coloring



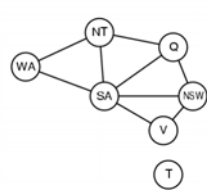
- **Solutions** are **complete** and **consistent** assignments (satisfying all constraints),
 - e.g., $WA = red, NT = green, Q = red, NSW = green, V = red, SA = blue, T = green$

B. Ommer | ommer@uni-heidelberg.de



Constraint graph

- **Binary CSP:** each constraint relates two variables
- **Constraint graph:** nodes are variables, arcs are constraints



- General-purpose CSP algorithms use the graph structure to speed up search. E.g., Tasmania is an independent subproblem!

B. Ommer | ommer@uni-heidelberg.de

Varieties of CSPs

- **Discrete variables**
 - finite domains:
 - n variables, domain size $d \rightarrow O(d^n)$ complete assignments
 - e.g., Boolean CSPs, incl. Boolean satisfiability (NP-complete)
 - infinite domains:
 - integers, strings, etc.
 - e.g., job scheduling, variables are start/end days for each job
 - need a constraint language, e.g., $StartJob_1 + 5 \leq StartJob_3$
 - linear constraints solvable, nonlinear undecidable
- **Continuous variables**
 - e.g., start/end times for Hubble Space Telescope observations
 - linear constraints solvable in polynomial time by linear programming

B. Ommer | ommer@uni-heidelberg.de

Varieties of constraints

- **Unary** constraints involve a single variable,
 - e.g., $SA \neq \text{green}$
- **Binary** constraints involve pairs of variables,
 - e.g., $SA \neq WA$
- **Higher-order** constraints involve 3 or more variables,
 - e.g., constraints from large neighborhoods
- **Preferences** (soft constraints), e.g., red is better than green often representable by a cost for each variable assignment
 - ⇒ **constrained optimization problems** (c.f. simplex)

B. Ommer | ommer@uni-heidelberg.de

Real-world CSPs

- **Assignment problems**
 - e.g., who teaches what class
- **Timetabling problems**
 - e.g., which class is offered when and where?
- **Transportation scheduling**
- **Factory scheduling**
- **Notice that many real-world problems involve real-valued variables**

B. Ommer | ommer@uni-heidelberg.de

Standard search formulation (incremental)

Let's start with the straightforward approach, then fix it

Idea from informed search (last chapter):

- Relaxed problems
- Optimal solution cost of a relaxed problem is not greater than the optimal solution cost of the real problem
- ⇨ cost of solving relaxed problem is an admissible heuristic for original problem
- Solve relax problem / then try to fix relaxed problem

B. Ommer | ommer@uni-heidelberg.de

Standard search formulation (incremental)

Let's start with the straightforward approach, then fix it

States are defined by the values assigned so far:

- **Initial state:** the empty assignment $\{\}$
- **Successor function:** assign a value to an unassigned variable that does not conflict with current assignment
 - fail if no legal assignments (not fixable)
- **Goal test:** the current assignment is complete

This is the same for all CSPs:

- Every solution appears at depth n with n variables
 - use depth-first search
- Path is irrelevant, so can also use complete-state formulation
- $b = (n - l) \cdot d$ at depth l , hence $n! \cdot d^n$ leaves!!

B. Ommer | ommer@uni-heidelberg.de

Backtracking search

- Variable assignments are **commutative**, i.e.,
 - [$WA = \text{red}$ then $NT = \text{green}$] same as [$NT = \text{green}$ then $WA = \text{red}$]

- Only need to consider assignments to a single variable at each node
 - simplifies to $b = d$ and number of leaves = d^n
- Depth-first search for CSPs with single-variable assignments is called **backtracking** search
- Backtracking search is the basic *uninformed* algorithm for CSPs
- Can solve n -queens for $n \approx 25$

Solutions for n queens problem:

n	1	2	3	4	5	6	7	8	9	10	...	24	25	26	27
fundamental:	1	0	0	1	2	1	6	12	46	92	...	28,439,272,956,934	275,906,683,743,434	2,789,712,466,510,289	29,363,791,967,678,199
all:	1	0	0	2	10	4	40	92	352	724	...	227,514,171,973,736	2,207,893,435,808,352	22,317,699,616,364,044	234,907,967,154,122,528

B. Ommer | ommer@uni-heidelberg.de

Backtracking search

```


function BACKTRACKING-SEARCH(csp) returns solution/failure
  return RECURSIVE-BACKTRACKING({ }, csp)
function RECURSIVE-BACKTRACKING(assignment, csp) returns soln/failure
  if assignment is complete then return assignment
  var ← SELECT-UNASSIGNED-VARIABLE(VARIABLES[csp], assignment, csp)
  for each value in ORDER-DOMAIN-VALUES(var, assignment, csp) do
    if value is consistent with assignment given CONSTRAINTS[csp] then
      add {var = value} to assignment
      result ← RECURSIVE-BACKTRACKING(assignment, csp)
      if result ≠ failure then return result
      remove {var = value} from assignment
  return failure

```


Backtracking example




Backtracking example



Backtracking example



Backtracking example



Improving backtracking efficiency

- General-purpose methods can give huge gains in speed:

- Which variable should be assigned next?
- In what order should its values be tried?
- Can we detect inevitable failure early?
- Can we take advantage of problem structure?

Most constrained variable

- **Most constrained variable:**
choose the variable with the fewest legal values




- a.k.a. **minimum remaining values (MRV)** heuristic

B. Ommer | ommer@uni-heidelberg.de

Most constraining variable

- **Tie-breaker among most constrained variables**
- **Most constraining variable:**
choose the variable with the most constraints on remaining variables



B. Ommer | ommer@uni-heidelberg.de

Least constraining value

- **Given a variable, choose the least constraining value:**
 - the one that rules out the fewest values in the remaining variables




- Combining these heuristics makes 1000 queens feasible

B. Ommer | ommer@uni-heidelberg.de

Forward checking

- **Idea:**
 - Keep track of **remaining legal values** for unassigned variables
 - Terminate search when any variable has no legal values



B. Ommer | ommer@uni-heidelberg.de

Forward checking


- **Idea:**
 - Keep track of remaining legal values for unassigned variables
 - Terminate search when any variable has no legal values



B. Ommer | ommer@uni-heidelberg.de

Forward checking


- **Idea:**
 - Keep track of remaining legal values for unassigned variables
 - Terminate search when any variable has no legal values



B. Ommer | ommer@uni-heidelberg.de

Forward checking


- Idea:
 - Keep track of remaining legal values for unassigned variables
 - Terminate search when any variable has no legal values



B. Ommer | ommer@uni-heidelberg.de

Constraint propagation

- Forward checking propagates information from assigned to unassigned variables, but doesn't provide early detection for all failures:




- NT and SA cannot both be blue!
- Constraint propagation repeatedly enforces constraints locally

B. Ommer | ommer@uni-heidelberg.de

Arc consistency


- Simplest form of propagation makes each arc consistent
- $X \rightarrow Y$ is consistent iff
for every value x of X there is some allowed y



B. Ommer | ommer@uni-heidelberg.de

Arc consistency


- Simplest form of propagation makes each arc consistent
- $X \rightarrow Y$ is consistent iff
for every value x of X there is some allowed y



B. Ommer | ommer@uni-heidelberg.de

Arc consistency

- Simplest form of propagation makes each arc consistent
- $X \rightarrow Y$ is consistent iff
for every value x of X there is some allowed y




- If X loses a value, neighbors of X need to be rechecked

B. Ommer | ommer@uni-heidelberg.de

Arc consistency

- Simplest form of propagation makes each arc consistent
- $X \rightarrow Y$ is consistent iff
for every value x of X there is some allowed y



- If X loses a value, neighbors of X need to be rechecked
- Arc consistency detects failure earlier than forward checking
- Can be run as a preprocessor or after each assignment

B. Ommer | ommer@uni-heidelberg.de

Arc consistency algorithm AC-3

```

function AC-3(csp) returns the CSP, possibly with reduced domains
inputs: csp, a binary CSP with variables  $\{X_1, X_2, \dots, X_n\}$ 
local variables: queue, a queue of arcs, initially all the arcs in csp
while queue is not empty do
  ( $X_i, X_j$ ) ← REMOVE-FIRST(queue)
  if REMOVE-INCONSISTENT-VALUES( $X_i, X_j$ ) then
    for each  $X_k$  in NEIGHBORS[ $X_i$ ] do
      add ( $X_k, X_i$ ) to queue


function REMOVE-INCONSISTENT-VALUES( $X_i, X_j$ ) returns true iff succeeds
removed ← false
for each  $x$  in DOMAIN[ $X_i$ ] do
  if no value  $y$  in DOMAIN[ $X_j$ ] allows  $(x, y)$  to satisfy the constraint  $X_i \leftrightarrow X_j$ 
  then delete  $x$  from DOMAIN[ $X_i$ ]; removed ← true
return removed

```

- Time complexity: $O(n^2d^3)$

B. Ommar | ommar@uni-heidelberg.de

Problem Structure



- Tasmania and mainland are independent subproblems
- Identifiable as connected components of constraint graph


B. Ommar | ommar@uni-heidelberg.de

Problem Structure

- Suppose each subproblem has c variables out of n total
- Worst-case solution cost is $n/c \cdot d^c$, linear in n
- E.g., $n=80, d=2, c=20$
 - $2^{80} = 4$ billion years at 10 million nodes/sec brute force vs.
 - $4 \cdot 2^{20} = 0.4$ seconds at 10 million nodes/sec

B. Ommar | ommar@uni-heidelberg.de

Tree-structured CSPs




- Theorem: if the constraint graph has no loops, the CSP can be solved in $O(n d^2)$ time
- Compare to general CSPs, where worst-case time is $O(d^n)$
- This property also applies to logical and probabilistic reasoning: an important example of the relation between syntactic restrictions and the complexity of reasoning.

B. Ommar | ommar@uni-heidelberg.de

Algorithm for tree-structured CSPs


- Choose a variable as root, order variables from root to leaves such that every node's parent precedes it in the ordering
- For j from n down to 2, apply RemoveInconsistent(Parent(X_j); X_j)
- For j from 1 to n , assign X_j consistently with Parent(X_j)



B. Ommar | ommar@uni-heidelberg.de

Nearly tree-structured CSPs

- Conditioning: instantiate a variable, prune its neighbors' domains: turn structure into tree



B. Ommar | ommar@uni-heidelberg.de

UNIVERSITÄT HEIDELBERG

Cutset Conditioning

- **Cutset conditioning:** instantiate (in all ways) a set of variables such that the remaining constraint graph is a tree
- Remove the values that conflict with the cutset assignment
- Solve the resulting tree structured CSPs

B. Ommer | ommer@uni-heidelberg.de

UNIVERSITÄT HEIDELBERG

Cutset Conditioning

- Cutset size $c \Rightarrow$ runtime $O(d^c (n - c)d^2)$, very fast for small c

B. Ommer | ommer@uni-heidelberg.de

UNIVERSITÄT HEIDELBERG

Local search for CSPs

- Hill-climbing, simulated annealing typically work with "complete" states, i.e., all variables assigned
- To apply to CSPs:
 - allow states with unsatisfied constraints
 - operators **reassign** variable values
- Variable selection: randomly select any conflicted variable
- Value selection by **min-conflicts** heuristic:
 - choose value that violates the fewest constraints
 - i.e., hill-climb with $h(n)$ = total number of violated constraints

B. Ommer | ommer@uni-heidelberg.de

UNIVERSITÄT HEIDELBERG

Example: 4-Queens

- **States:** 4 queens in 4 columns ($4^4 = 256$ states)
- **Actions:** move queen in column
- **Goal test:** no attacks
- **Evaluation:** $h(n)$ = number of attacks

- Given random initial state, can solve n -queens in almost constant time for arbitrary n with high probability (e.g., $n = 10,000,000$)

B. Ommer | ommer@uni-heidelberg.de

UNIVERSITÄT HEIDELBERG

Performance of min-conflicts

- Given random initial state, can solve n -queens in almost constant time for arbitrary n with high probability (e.g., $n = 10,000,000$)
- The same appears to be true for any randomly-generated CSP except in a narrow range of the ratio

$$R = \frac{\text{number of constraints}}{\text{number of variables}}$$

B. Ommer | ommer@uni-heidelberg.de

UNIVERSITÄT HEIDELBERG

Summary

- CSPs are a special kind of problem:
 - states defined by values of a fixed set of variables
 - goal test defined by constraints on variable values
- Backtracking = depth-first search with one variable assigned per node
- Variable ordering and value selection heuristics help significantly
- Forward checking prevents assignments that guarantee later failure
- Constraint propagation (e.g., arc consistency) does additional work to constrain values and detect inconsistencies
- Iterative min-conflicts is usually effective in practice

B. Ommer | ommer@uni-heidelberg.de