



Artificial Intelligence
Prof. Björn Ommer
HCI & IWR






Nov 21, 2017 — Adversarial Search & Games

UNIVERSITÄT HEIDELBERG

Outline – Game Playing & Adversarial Search

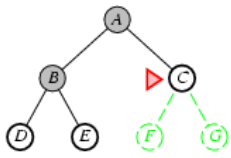
- Games
- Perfect play
 - minimax decisions
 - α - β pruning
- Resource limits and approximate evaluation
- Games of chance
- Games of imperfect information

B. Ommer | ommer@uni-heidelberg.de 2

UNIVERSITÄT HEIDELBERG

Games vs. search problems

- "Unpredictable" opponent \rightarrow specifying a move for every possible opponent reply



- Time limits \rightarrow unlikely to find goal, must approximate

B. Ommer | ommer@uni-heidelberg.de 3

UNIVERSITÄT HEIDELBERG

Games vs. search problems

- Brute-force, adding pruning strategies, ... automatically learning evaluation strategies:
 - Computer considers possible lines of play (Babbage, 1846)
 - Algorithm for perfect play (Zermelo, 1912; Von Neumann, 1944)
 - Finite horizon, approximate evaluation (Zuse, 1945; Wiener, 1948; Shannon, 1950)
 - First chess program (Turing, 1951)
 - Machine learning to improve evaluation accuracy (Samuel, 1952-57)
 - Pruning to allow deeper search (McCarthy, 1956)
 - ...
 - DeepBlue beating world champion in chess, 1997
 - Deep Learning: AlphaGo beating champion in Go, 2016

B. Ommer | ommer@uni-heidelberg.de 4

UNIVERSITÄT HEIDELBERG

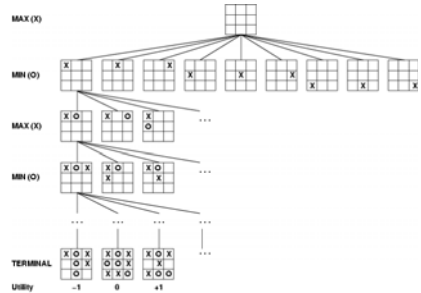
Types of games

	deterministic	chance
perfect information	chess, checkers, go, othello	backgammon monopoly
imperfect information	battleships, blind tictactoe	bridge, poker, scrabble nuclear war

B. Ommer | ommer@uni-heidelberg.de 5

UNIVERSITÄT HEIDELBERG


Game tree (2-player, deterministic, turns)



B. Ommer | ommer@uni-heidelberg.de 6

Minimax

- Perfect play for deterministic games
- Idea: choose move to position with highest **minimax value**
= best achievable payoff against best play
- E.g., 2-ply game:



B. Ommer | ommer@uni-heidelberg.de

7

Minimax algorithm

```
function MINIMAX-DECISION(state) returns an action
  inputs: state, current state in game
  return the a in ACTIONS(state) maximizing MIN-VALUE(RESULT(a, state))
```

```
function MAX-VALUE(state) returns a utility value
  if TERMINAL-TEST(state) then return UTILITY(state)
   $v \leftarrow -\infty$ 
  for a, s in SUCCESSORS(state) do  $v \leftarrow \text{MAX}(v, \text{MIN-VALUE}(s))$ 
  return v
```

```
function MIN-VALUE(state) returns a utility value
  if TERMINAL-TEST(state) then return UTILITY(state)
   $v \leftarrow \infty$ 
  for a, s in SUCCESSORS(state) do  $v \leftarrow \text{MIN}(v, \text{MAX-VALUE}(s))$ 
  return v
```

B. Ommer | ommer@uni-heidelberg.de

8


Properties of minimax

- Complete?** Yes (if tree is finite)
- Optimal?** Yes (against an optimal opponent)
- Time complexity?** $O(b^m)$
- Space complexity?** $O(bm)$ (depth-first exploration)
- For chess, $b \approx 35$, $m \approx 100$ for "reasonable" games
→ exact solution completely infeasible (35^{100})

B. Ommer | ommer@uni-heidelberg.de

9


α - β pruning example



B. Ommer | ommer@uni-heidelberg.de

10


α - β pruning example



B. Ommer | ommer@uni-heidelberg.de

11


α - β pruning example




B. Ommer | ommer@uni-heidelberg.de

12

α - β pruning example



α - β pruning example




Properties of α - β

- Pruning **does not** affect final result
- Good move ordering improves effectiveness of pruning
- With "perfect ordering," time complexity = $O(b^{m/2})$
→ **doubles** depth of search
- A simple example of the value of reasoning about which computations are relevant (a form of **metareasoning**)

Why is it called α - β ?

- α is the value of the best (i.e., highest-value) choice found so far at any choice point along the path for *max*
- If v is worse than α , *max* will avoid it → prune that branch
- Define β similarly for *min*



The α - β algorithm

```

function ALPHA-BETA-SEARCH(state) returns an action
  inputs: state, current state in game
   $v \leftarrow \text{MAX-VALUE}(\text{state}, -\infty, +\infty)$ 
  return the action in  $\text{SUCCESSORS}(\text{state})$  with value  $v$ 

function MAX-VALUE(state,  $\alpha$ ,  $\beta$ ) returns a utility value
  inputs: state, current state in game
          $\alpha$ , the value of the best alternative for MAX along the path to state
          $\beta$ , the value of the best alternative for MIN along the path to state
  if  $\text{TERMINAL-TEST}(\text{state})$  then return  $\text{UTILITY}(\text{state})$ 
   $v \leftarrow -\infty$ 
  for  $a, s$  in  $\text{SUCCESSORS}(\text{state})$  do
     $v \leftarrow \text{MAX}(v, \text{MIN-VALUE}(s, \alpha, \beta))$ 
    if  $v \geq \beta$  then return  $v$ 
     $\alpha \leftarrow \text{MAX}(\alpha, v)$ 
  return  $v$ 

```

The α - β algorithm

```

function MIN-VALUE(state,  $\alpha$ ,  $\beta$ ) returns a utility value
  inputs: state, current state in game
          $\alpha$ , the value of the best alternative for MAX along the path to state
          $\beta$ , the value of the best alternative for MIN along the path to state
  if  $\text{TERMINAL-TEST}(\text{state})$  then return  $\text{UTILITY}(\text{state})$ 
   $v \leftarrow +\infty$ 
  for  $a, s$  in  $\text{SUCCESSORS}(\text{state})$  do
     $v \leftarrow \text{MIN}(v, \text{MAX-VALUE}(s, \alpha, \beta))$ 
    if  $v \leq \alpha$  then return  $v$ 
     $\beta \leftarrow \text{MIN}(\beta, v)$ 
  return  $v$ 

```

Resource limits

Suppose we have 100 secs, explore 10^4 nodes/sec

→ 10^6 nodes per move $\sim 35^{8/2}$

→ α - β reaches depth 8 → pretty good chess program


Standard approach:

- **cutoff test (instead of terminal test):**
e.g., depth limit (perhaps add **quiescence search**)
- **evaluation function (instead of utility):**
= estimated desirability of position

B. Ommer | ommer@uni-heidelberg.de

20

Evaluation functions



Black to move

White slightly better

White to move


Black winning

- For chess, typically **linear** weighted sum of **features**
 $Eval(s) = w_1 f_1(s) + w_2 f_2(s) + \dots + w_n f_n(s)$
- e.g., $w_1 = 9$ with
 $f_1(s) = (\text{number of white queens}) - (\text{number of black queens}), \text{ etc.}$

B. Ommer | ommer@uni-heidelberg.de

21

(Digression: Exact values don't matter)



- Behavior is preserved under any monotonic transformation of Eval
- Only the order matters:
 - payoff in deterministic games acts as an ordinal utility function

B. Ommer | ommer@uni-heidelberg.de

22

Cutting off search

MinimaxCutoff is identical to **MinimaxValue** except

1. **Terminal?** is replaced by **Cutoff?**
2. **Utility** is replaced by **Eval**

Does it work in practice?

$$b^m = 10^6, b=35 \rightarrow m=4$$

4-ply lookahead is a hopeless chess player!

- 4-ply \approx human novice
- 8-ply \approx typical PC, human master
- 12-ply \approx Deep Blue, Kasparov

B. Ommer | ommer@uni-heidelberg.de

23


Deterministic games in practice

- Checkers: Chinook ended 40-year-reign of human world champion Marion Tinsley in 1994. Used a precomputed endgame database defining perfect play for all positions involving 8 or fewer pieces on the board, a total of 444 billion positions.
- Chess: Deep Blue defeated human world champion Garry Kasparov in a six-game match in 1997. Deep Blue searches 200 million positions per second, uses very sophisticated evaluation, and undisclosed methods for extending some lines of search up to 40 ply.
- Othello: human champions refuse to compete against computers, who are *too good*.
- *Until 2016*: Go: human champions refuse to compete against computers, who are *too bad*. In go, $b > 300$, so most programs use pattern knowledge bases to suggest plausible moves. *However, see later slide...*

B. Ommer | ommer@uni-heidelberg.de

24

Nondeterministic games: backgammon




B. Ommer | ommer@uni-heidelberg.de

25

Nondeterministic games in general

- In nondeterministic games, chance introduced by dice, card-shuffling
- Simplified example with coin-flipping:



Algorithm for nondeterministic games

- Expectiminimax gives perfect play
- Just like Minimax, except we must also handle chance nodes:

```


...
if state is a MAX node then
    return the highest EXPECTIMINIMAX-VALUE of SUCCESSORS(state)
if state is a MIN node then
    return the lowest EXPECTIMINIMAX-VALUE of SUCCESSORS(state)
if state is a chance node then
    return average of EXPECTIMINIMAX-VALUE of SUCCESSORS(state)
...

```

Nondeterministic games in practice

- Dice rolls **increase b**: 21 possible rolls with 2 dice
- Backgammon 20 legal moves (can be 6,000 with 1-1 roll)
- As depth increases, probability of reaching a given node shrinks
 - ⇒ value of lookahead is diminished
- α - β pruning is much less effective
- TDGammon uses depth-2 search + very good Eval ⇒ world-champion level

Digression: Exact values DO matter



- Behavior is preserved only by positive linear transformation of Eval
- Hence Eval should be proportional to the expected payoff

Games of imperfect information

- E.g., card games, where opponent's initial cards are unknown
- Typically we can calculate a probability for each possible deal
- Seems just like having one big dice roll at the beginning of the game*
- Idea: compute the minimax value of each action in each deal, then choose the action with **highest expected value over all deals***
- Special case: if an action is optimal for all deals, it's optimal.*
- GIB (best bridge program for long) approximates this idea by
 - generating 100 deals consistent with bidding information
 - picking the action that wins most tricks on average

Commonsense example

- Road A leads to a small heap of gold pieces
- Road B leads to a fork:
 - take the left fork and you'll find a mound of jewels;
 - take the right fork and you'll be run over by a bus.

Commonsense example

- Road A leads to a small heap of gold pieces
- Road B leads to a fork:
 - take the left fork and you'll find a mound of jewels;
 - take the right fork and you'll be run over by a bus.
- Road A leads to a small heap of gold pieces
- Road B leads to a fork:
 - take the left fork and you'll be run over by a bus;
 - take the right fork and you'll find a mound of jewels.

Commonsense example

- Road A leads to a small heap of gold pieces
- Road B leads to a fork:
 - take the left fork and you'll find a mound of jewels;
 - take the right fork and you'll be run over by a bus.
- Road A leads to a small heap of gold pieces
- Road B leads to a fork:
 - take the left fork and you'll be run over by a bus;
 - take the right fork and you'll find a mound of jewels.
- Road A leads to a small heap of gold pieces
- Road B leads to a fork:
 - guess correctly and you'll find a mound of jewels;
 - guess incorrectly and you'll be run over by a bus.

Proper Analysis

- *** Intuition that the value of an action is the average of its values in all actual states is WRONG**
- With partial observability, value of an action depends on the information state or belief state the agent is in
- Can generate and search a tree of information states
- Leads to rational behaviors such as
 - Acting to obtain information, exploration, ...
 - Signalling to one's partner
 - Acting randomly to minimize information disclosure

Summary

- Games are fun to work on!
- They illustrate several important points about AI
 - perfection is unattainable → must approximate
 - good idea to think about what to think about
 - uncertainty constrains the assignment of values to states
 - optimal decisions depend on information state, not real state
- Games are to AI as grand prix racing is to automobile design




- ... recent improvements... AlphaGo zero



Intelligent Game Playing

- Learning optimal actions


 x

 $y=f(x;w)$

- Learn a function f with some parameters w to predict next move y
- ⇒ Machine learning and esp. deep learning